



Sharif-OS-Lab /  
summer1403-5-99101087\_99100422 🔔

[Code](#)[Issues](#) **2**[Pull requests](#)[Actions](#)[Projects](#)[Security](#)[Insights](#)[Edit](#)[New issue](#)[Jump to bottom](#)

# Session 5 Report #1

[Open](#)[22 tasks done](#)

Amirreza81 opened this issue 5 days ago · 0 comments

Assignees



Labels

[documentation](#)

Amirreza81 commented 5 days ago • edited ▾

Team Name: 99101087-99100422

Student Name of member 1: AmirReza Azari

Student No. of member 1: 99101087

Student Name of member 2: Bozorgmehr Zia

Student No. of member 2: 99100422

☒ Read Session Contents.

## Section 5.3.1

☒ Write the Hello World! program☒ FILL HERE with your source code

The program is here in Hello\_World.c:

```
GNU nano 2.2.6                               File: Hello_World.c

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/wait.h>

int main() {
    int pipes[2];
    if (pipe(pipes) < 0) {
        perror("Error!");
        exit(1);
    }
    char buffer[32];
    if (fork() == 0) {
        read(pipes[0], buffer, sizeof(buffer));
        puts(buffer);
    } else {
        strcpy(buffer, "Hello World!");
        write(pipes[1], buffer, sizeof(buffer));
        wait(NULL);
    }
}
```

Then we have:

```
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <sys/wait.h>

int main() {
    int pipes[2];
    if (pipe(pipes) < 0) {
        perror("Error!");
        exit(1);
    }
    char buffer[32];
    if (fork() == 0) {
        read(pipes[0], buffer, sizeof(buffer));
        puts(buffer);
    } else {
        strcpy(buffer, "Hello World!");
        write(pipes[1], buffer, sizeof(buffer));
        wait(NULL);
    }
}
```

[ Wrote 22 lines ]

```
root@debian:~# gcc -o 1_out Hello_World.c
root@debian:~# ./1_out
Hello World!
root@debian:~#
```

☒ Write the `ls` to `wc` program

☒ FILL HERE with your source code

The program is here in `ls_to_wc.c`:

```
GNU nano 2.2.6                               File: ls_to_wc.c

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main() {
    int pipes[2];
    if (pipe(pipes) < 0) {
        perror("Error!");
        exit(1);
    }
    if (fork() == 0) {
        close(pipes[1]); // do not need the writing part
        dup2(pipes[0], STDIN_FILENO); // change to be stdin
        close(pipes[0]);
        execlp("wc", "wc", NULL);
    } else {
        close(pipes[0]); // do not need the reading part
        dup2(pipes[1], STDOUT_FILENO); // change to be stdout
        close(pipes[1]);
        execlp("ls", "ls", NULL);
    }
}
```

Then the result is:

```

int pipes[2];
if (pipe(pipes) < 0) {
    perror("Error!");
    exit(1);
}
if (fork() == 0) {
    close(pipes[1]); // do not need the writing part
    dup2(pipes[0], STDIN_FILENO); // change to be stdin
    close(pipes[0]);
    execlp("wc", "wc", NULL);
} else {
    close(pipes[0]); // do not need the reading part
    dup2(pipes[1], STDOUT_FILENO); // change to be stdout
    close(pipes[1]);
    execlp("ls", "ls", NULL);
}
}

```

[ Wrote 22 lines ]

```

root@debian:~# make ls_to_wc
cc      ls_to_wc.c  -o ls_to_wc
root@debian:~# ./ls_to_wc
root@debian:~#          4          4          40

root@debian:~# ls|wc -l
4
root@debian:~#

```

### ☑ Investigate how to have a bi-direction pipe

#### ☑ FILL HERE with your descriptions

We have two options.

- i. We can either just use two pipes (one from parent to child and one vice versa)
- ii. We can use unix sockets. Unix sockets are somehow a combination of TCP/UDP sockets and named pipes. They can operate in three different ways. SOCK\_STREAM , SOCK\_DGRAM and SOCK\_SEQPACKET.

For an experience like pipes, we can use SOCK\_DGRAM which is mostly like UDP packets. Unix sockets can improve the speed of TCP and UDP connections in compared to sending packets to loopback. Thus it's recommended to use unix sockets for reverse proxying.

## Section 5.3.2

First we do `man 7 signal`. The result is:

```
SIGNAL(7)                                Linux Programmer's Manual                                SIGNAL(7)

NAME
    signal - overview of signals

DESCRIPTION
    Linux supports both POSIX reliable signals (hereinafter "standard signals") and POSIX
    real-time signals.

    Signal dispositions
    Each signal has a current disposition, which determines how the process behaves when it is
    delivered the signal.

    The entries in the "Action" column of the tables below specify the default disposition for
    each signal, as follows:

    Term    Default action is to terminate the process.

    Ign     Default action is to ignore the signal.

    Core    Default action is to terminate the process and dump core (see core(5)).

    Stop    Default action is to stop the process.

    Cont    Default action is to continue the process if it is currently stopped.

    A process can change the disposition of a signal using sigaction(2) or signal(2). (The
    latter is less portable when establishing a signal handler; see signal(2) for details.)
    Using these system calls, a process can elect one of the following behaviors to occur on
    delivery of the signal: perform the default action; ignore the signal; or catch the signal
    with a signal handler, a programmer-defined function that is automatically invoked when
    the signal is delivered. (By default, the signal handler is invoked on the normal process
    stack. It is possible to arrange that the signal handler uses an alternate stack; see
    sigaltstack(2) for a discussion of how to do this and when it might be useful.)

    The signal disposition is a per-process attribute: in a multithreaded application, the
    Manual page signal(7) line 1 (press h for help or q to quit)
```

A process can change the disposition of a signal using `sigaction(2)` or `signal(2)`. (The latter is less portable when establishing a signal handler; see `signal(2)` for details.) Using these system calls, a process can elect one of the following behaviors to occur on delivery of the signal: perform the default action; ignore the signal; or catch the signal with a `signal handler`, a programmer-defined function that is automatically invoked when the signal is delivered. (By default, the signal handler is invoked on the normal process stack. It is possible to arrange that the signal handler uses an alternate stack; see `sigaltstack(2)` for a discussion of how to do this and when it might be useful.)

The signal disposition is a per-process attribute: in a multithreaded application, the disposition of a particular signal is the same for all threads.

A child created via `fork(2)` inherits a copy of its parent's signal dispositions. During an `execve(2)`, the dispositions of handled signals are reset to the default; the dispositions of ignored signals are left unchanged.

#### Sending a signal

The following system calls and library functions allow the caller to send a signal:

<code>raise(3)</code>	Sends a signal to the calling thread.
<code>kill(2)</code>	Sends a signal to a specified process, to all members of a specified process group, or to all processes on the system.
<code>killpg(2)</code>	Sends a signal to all of the members of a specified process group.
<code>pthread_kill(3)</code>	Sends a signal to a specified POSIX thread in the same process as the caller.
<code>tgkill(2)</code>	Sends a signal to a specified thread within a specific process. (This is the system call used to implement <code>pthread_kill(3)</code> .)
<code>sigqueue(3)</code>	Sends a real-time signal with accompanying data to a specified process.

#### Waiting for a signal to be caught

Manual page `signal(7)` line 28 (press h for help or q to quit)

### Waiting for a signal to be caught

The following system calls suspend execution of the calling process or thread until a signal is caught (or an unhandled signal terminates the process):

`pause(2)` Suspend execution until any signal is caught.

`sigsuspend(2)` Temporarily changes the signal mask (see below) and suspends execution until one of the unmasked signals is caught.

### Synchronously accepting a signal

Rather than asynchronously catching a signal via a signal handler, it is possible to synchronously accept the signal, that is, to block execution until the signal is delivered, at which point the kernel returns information about the signal to the caller. There are two general ways to do this:

- \* `sigwaitinfo(2)`, `sigtimedwait(2)`, and `sigwait(3)` suspend execution until one of the signals in a specified set is delivered. Each of these calls returns information about the delivered signal.

- \* `signalfd(2)` returns a file descriptor that can be used to read information about signals that are delivered to the caller. Each `read(2)` from this file descriptor blocks until one of the signals in the set specified in the `signalfd(2)` call is delivered to the caller. The buffer returned by `read(2)` contains a structure describing the signal.

### Signal mask and pending signals

A signal may be **blocked**, which means that it will not be delivered until it is later unblocked. Between the time when it is generated and when it is delivered a signal is said to be **pending**.

Each thread in a process has an independent **signal mask**, which indicates the set of signals that the thread is currently blocking. A thread can manipulate its signal mask using `pthread_sigmask(3)`. In a traditional single-threaded application, `sigprocmask(2)` can be used to manipulate the signal mask.

A child created via `fork(2)` inherits a copy of its parent's signal mask; the signal mask

Manual page `signal(7)` line 62 (press h for help or q to quit)



is preserved across `execve(2)`.

A signal may be generated (and thus pending) for a process as a whole (e.g., when sent using `kill(2)`) or for a specific thread (e.g., certain signals, such as `SIGSEGV` and `SIGFPE`, generated as a consequence of executing a specific machine-language instruction are thread directed, as are signals targeted at a specific thread using `pthread_kill(3)`). A process-directed signal may be delivered to any one of the threads that does not currently have the signal blocked. If more than one of the threads has the signal unblocked, then the kernel chooses an arbitrary thread to which to deliver the signal.

A thread can obtain the set of signals that it currently has pending using `sigpending(2)`. This set will consist of the union of the set of pending process-directed signals and the set of signals pending for the calling thread.

A child created via `fork(2)` initially has an empty pending signal set; the pending signal set is preserved across an `execve(2)`.

#### Standard signals

Linux supports the standard signals listed below. Several signal numbers are architecture-dependent, as indicated in the "Value" column. (Where three values are given, the first one is usually valid for alpha and sparc, the middle one for x86, arm, and most other architectures, and the last one for mips. (Values for parisc are *not* shown; see the Linux kernel source for signal numbering on that architecture.) A - denotes that a signal is absent on the corresponding architecture.)

First the signals described in the original POSIX.1-1990 standard.

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from <code>abort(3)</code>
Manual page <code>signal(7)</code> line 98 (press h for help or q to quit)			

First the signals described in the original POSIX.1-1990 standard.

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from abort(3)
SIGFPE	8	Core	Floating point exception
SIGKILL	9	Term	Kill signal
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from alarm(2)
SIGTERM	15	Term	Termination signal
SIGUSR1	30,10,16	Term	User-defined signal 1
SIGUSR2	31,12,17	Term	User-defined signal 2
SIGCHLD	20,17,18	Ign	Child stopped or terminated
SIGCONT	19,18,25	Cont	Continue if stopped
SIGSTOP	17,19,23	Stop	Stop process
SIGTSTP	18,20,24	Stop	Stop typed at terminal
SIGTTIN	21,21,26	Stop	Terminal input for background process
SIGTTOU	22,22,27	Stop	Terminal output for background process

The signals SIGKILL and SIGSTOP cannot be caught, blocked, or ignored.

Next the signals not in the POSIX.1-1990 standard but described in SUSv2 and POSIX.1-2001.

Signal	Value	Action	Comment
SIGBUS	10,7,10	Core	Bus error (bad memory access)
SIGPOLL		Term	Pollable event (Sys V). Synonym for SIGIO

Manual page signal(7) line 123 (press h for help or q to quit)

Signal	Value	Action	Comment
SIGBUS	10,7,10	Core	Bus error (bad memory access)
SIGPOLL		Term	Pollable event (Sys V). Synonym for SIGIO
SIGPROF	27,27,29	Term	Profiling timer expired
SIGSYS	12,31,12	Core	Bad argument to routine (SVr4)
SIGTRAP	5	Core	Trace/breakpoint trap
SIGURG	16,23,21	Ign	Urgent condition on socket (4.2BSD)
SIGVTALRM	26,26,28	Term	Virtual alarm clock (4.2BSD)
SIGXCPU	24,24,30	Core	CPU time limit exceeded (4.2BSD)
SIGXFSZ	25,25,31	Core	File size limit exceeded (4.2BSD)

Up to and including Linux 2.2, the default behavior for SIGSYS, SIGXCPU, SIGXFSZ, and (on architectures other than SPARC and MIPS) SIGBUS was to terminate the process (without a core dump). (On some other UNIX systems the default action for SIGXCPU and SIGXFSZ is to terminate the process without a core dump.) Linux 2.4 conforms to the POSIX.1-2001 requirements for these signals, terminating the process with a core dump.

Next various other signals.

Signal	Value	Action	Comment
SIGIOT	6	Core	IOT trap. A synonym for SIGABRT
SIGEMT	7,-,7	Term	
SIGSTKFLT	-,16,-	Term	Stack fault on coprocessor (unused)
SIGIO	23,29,22	Term	I/O now possible (4.2BSD)
SIGCLD	-,-,18	Ign	A synonym for SIGCHLD
SIGPWR	29,30,19	Term	Power failure (System V)
SIGINFO	29,-,-		A synonym for SIGPWR
SIGLOST	-,,-	Term	File lock lost (unused)
SIGWINCH	28,28,20	Ign	Window resize signal (4.3BSD, Sun)
SIGUNUSED	-,31,-	Core	Synonymous with SIGSYS

(Signal 29 is SIGINFO / SIGPWR on an alpha but SIGLOST on a sparc.)

Manual page signal(7) line 153 (press h for help or q to quit)

`SIGEMT` is not specified in POSIX.1-2001, but nevertheless appears on most other UNIX systems, where its default action is typically to terminate the process with a core dump.

`SIGPWR` (which is not specified in POSIX.1-2001) is typically ignored by default on those other UNIX systems where it appears.

`SIGIO` (which is not specified in POSIX.1-2001) is ignored by default on several other UNIX systems.

Where defined, `SIGUNUSED` is synonymous with `SIGSYS` on most architectures.

#### Real-time signals

Linux supports real-time signals as originally defined in the POSIX.1b real-time extensions (and now included in POSIX.1-2001). The range of supported real-time signals is defined by the macros `SIGRTMIN` and `SIGRTMAX`. POSIX.1-2001 requires that an implementation support at least `_POSIX_RTSIG_MAX` (8) real-time signals.

The Linux kernel supports a range of 32 different real-time signals, numbered 33 to 64. However, the glibc POSIX threads implementation internally uses two (for NPTL) or three (for LinuxThreads) real-time signals (see `pthread(7)`), and adjusts the value of `SIGRTMIN` suitably (to 34 or 35). Because the range of available real-time signals varies according to the glibc threading implementation (and this variation can occur at run time according to the available kernel and glibc), and indeed the range of real-time signals varies across UNIX systems, programs should **never refer to real-time signals using hard-coded numbers**, but instead should always refer to real-time signals using the notation `SIGRTMIN+n`, and include suitable (run-time) checks that `SIGRTMIN+n` does not exceed `SIGRTMAX`.

Unlike standard signals, real-time signals have no predefined meanings: the entire set of real-time signals can be used for application-defined purposes.

The default action for an unhandled real-time signal is to terminate the receiving process.

Real-time signals are distinguished by the following:

```
Manual page signal(7) line 189 (press h for help or q to quit)
```

- ☒ Describe the usecase of different signals:

✓ [FILL HERE with the description for SIGINT.]

**SIGINT** : SIGINT is the signal sent when we press Ctrl+C. The default action is to terminate the process. However, some programs override this action and handle it differently. So, it is a signal number 2. Is send when we press ctrl + c in a running program. (Interrupt from keyboard).

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
<b>SIGINT</b>	<b>2</b>	<b>Term</b>	<b>Interrupt from keyboard</b>
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from <code>abort(3)</code>
SIGFPE	8	Core	Floating point exception
SIGKILL	9	Term	Kill signal
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from <code>alarm(2)</code>
SIGTERM	15	Term	Termination signal
SIGUSR1	30,10,16	Term	User-defined signal 1
SIGUSR2	31,12,17	Term	User-defined signal 2
SIGCHLD	20,17,18	Ign	Child stopped or terminated
SIGCONT	19,18,25	Cont	Continue if stopped
SIGSTOP	17,19,23	Stop	Stop process
SIGTSTP	18,20,24	Stop	Stop typed at terminal
SIGTTIN	21,21,26	Stop	Terminal input for background process
SIGTTOU	22,22,27	Stop	Terminal output for background process

The signals **SIGKILL** and **SIGSTOP** cannot be caught, blocked, or ignored.

✓ [FILL HERE with the description for SIGHUP.]

**SIGHUP** : ("signal hang up") is a signal sent to a process when its controlling terminal is closed. It was originally designed to notify the process of a serial line drop. SIGHUP is a symbolic constant defined in the header file signal.h.

In other words, the SIGHUP signal was sent to a process when its controlling terminal is closed. Nowadays, it's often used to indicate that a process should reload its configuration or restart. The default action of this signal is to terminate the process.

Signal	Value	Action	Comment
<b>SIGHUP</b>	<b>1</b>	<b>Term</b>	<b>Hangup detected on controlling terminal or death of controlling process</b>
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from <b>abort(3)</b>
SIGFPE	8	Core	Floating point exception
SIGKILL	9	Term	Kill signal
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from <b>alarm(2)</b>
SIGTERM	15	Term	Termination signal
SIGUSR1	30,10,16	Term	User-defined signal 1
SIGUSR2	31,12,17	Term	User-defined signal 2
SIGCHLD	20,17,18	Ign	Child stopped or terminated
SIGCONT	19,18,25	Cont	Continue if stopped
SIGSTOP	17,19,23	Stop	Stop process
SIGTSTP	18,20,24	Stop	Stop typed at terminal
SIGTTIN	21,21,26	Stop	Terminal input for background process
SIGTTOU	22,22,27	Stop	Terminal output for background process

The signals **SIGKILL** and **SIGSTOP** cannot be caught, blocked, or ignored.

✓ [FILL HERE with the description for SIGSTOP.]

**SIGSTOP** : The SIGSTOP signal instructs a process to halt its execution temporarily. Unlike SIGINT, SIGSTOP cannot be caught, blocked, or ignored by the process. The process will remain stopped until it receives a SIGCONT signal.

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from <code>abort(3)</code>
SIGFPE	8	Core	Floating point exception
SIGKILL	9	Term	Kill signal
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from <code>alarm(2)</code>
SIGTERM	15	Term	Termination signal
SIGUSR1	30,10,16	Term	User-defined signal 1
SIGUSR2	31,12,17	Term	User-defined signal 2
SIGCHLD	20,17,18	Ign	Child stopped or terminated
SIGCONT	19,18,25	Cont	Continue if stopped
<b>SIGSTOP</b>	<b>17,19,23</b>	<b>Stop</b>	<b>Stop process</b>
SIGTSTP	18,20,24	Stop	Stop typed at terminal
SIGTTIN	21,21,26	Stop	Terminal input for background process
SIGTTOU	22,22,27	Stop	Terminal output for background process

The signals SIGKILL and **SIGSTOP** cannot be caught, blocked, or ignored.

✓ [FILL HERE with the description for SIGCONT.]

**SIGCONT** : The SIGCONT signal is sent to a process to instruct it to continue execution if it was previously stopped by SIGSTOP or SIGTSTP. It's commonly used to resume processes that were paused or suspended.

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from <code>abort(3)</code>
SIGFPE	8	Core	Floating point exception
SIGKILL	9	Term	Kill signal
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from <code>alarm(2)</code>
SIGTERM	15	Term	Termination signal
SIGUSR1	30,10,16	Term	User-defined signal 1
SIGUSR2	31,12,17	Term	User-defined signal 2
SIGCHLD	20,17,18	Ign	Child stopped or terminated
SIGCONT	19,18,25	Cont	Continue if stopped
SIGSTOP	17,19,23	Stop	Stop process
SIGTSTP	18,20,24	Stop	Stop typed at terminal
SIGTTIN	21,21,26	Stop	Terminal input for background process
SIGTTOU	22,22,27	Stop	Terminal output for background process

The signals SIGKILL and SIGSTOP cannot be caught, blocked, or ignored.

Summary:

Wikipedia has great definitions for SIGSTOP:

*“When SIGSTOP is sent to a process, the usual behaviour is to pause that process in its current state. The process will only resume execution if it is sent the SIGCONT signal. SIGSTOP and SIGCONT are used for job control in the Unix shell, among other purposes. SIGSTOP cannot be caught or ignored.”*

and SIGCONT:

*“When SIGSTOP or SIGTSTP is sent to a process, the usual behaviour is to pause that process in its current state. The process will only resume execution if it is sent the SIGCONT signal. SIGSTOP and SIGCONT are used for job control in the Unix shell, among other purposes.”*

In short, SIGSTOP tells a process to “hold on” and SIGCONT tells a process to “pick up where you left off”. This can work really well for rsync jobs since you can pause the job, clear up some space on the destination device, and then resume the job. The source rsync process just thinks that the destination rsync process is taking a long time to respond.



✓ [FILL HERE with the description for SIGKILL.]

**SIGKILL** : Kill signal. This signal cannot be caught, blocked, or ignored. It's often used as a last resort to terminate

unresponsive or misbehaving processes that cannot be terminated through other means.

The signals **SIGKILL** and **SIGSTOP** cannot be caught, blocked, or ignored.

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from <code>abort(3)</code>
SIGFPE	8	Core	Floating point exception
<b>SIGKILL</b>	<b>9</b>	<b>Term</b>	<b>Kill signal</b>
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from <code>alarm(2)</code>
SIGTERM	15	Term	Termination signal
SIGUSR1	30,10,16	Term	User-defined signal 1
SIGUSR2	31,12,17	Term	User-defined signal 2
SIGCHLD	20,17,18	Ign	Child stopped or terminated
SIGCONT	19,18,25	Cont	Continue if stopped
SIGSTOP	17,19,23	Stop	Stop process
SIGTSTP	18,20,24	Stop	Stop typed at terminal
SIGTTIN	21,21,26	Stop	Terminal input for background process
SIGTTOU	22,22,27	Stop	Terminal output for background process

The signals **SIGKILL** and **SIGSTOP** cannot be caught, blocked, or ignored.

## What is SIGKILL (signal 9)

**SIGKILL** is a type of communication, known as a signal, used in Unix or Unix-like operating systems like Linux to immediately terminate a process. It is used by Linux operators, and also by container orchestrators like Kubernetes, when they need to shut down a container or pod on a Unix-based operating system.

A signal is a standardized message sent to a running program that triggers a specific action (such as terminating or handling an error). It is a type of Inter Process Communication (IPC). When an operating system sends a signal to a target process, it waits for atomic instructions to complete, and then interrupts the execution of the process, and handles the signal.

**SIGKILL** instructs the process to terminate immediately. It cannot be ignored or blocked. The process is killed, and if it is running threads, those are killed as well. If the **SIGKILL** signal fails to terminate a process and its threads, this indicates an operating system malfunction.

✓ Describe **SIGALRM**

☑ [FILL HERE with your description.]

**SIGALRM** : SIGALRM is an asynchronous signal. The SIGALRM signal is raised when a time interval specified in a call to the alarm or alarmd function expires. Because SIGALRM is an asynchronous signal, the SAS/C library discovers the signal only when you call a function, when a function returns, or when you issue a call to sigchk. In other words, SIGALRM is a signal sent to a process when the real-time clock timer set by alarm() or setitimer() expires. It's often used in programming to handle timeouts or to perform periodic tasks. When a process receives SIGALRM, it typically interrupts its current execution and invokes a signal handler, if one has been registered for this signal.

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from abort(3)
SIGFPE	8	Core	Floating point exception
SIGKILL	9	Term	Kill signal
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from alarm(2)
SIGTERM	15	Term	Termination signal
SIGUSR1	30,10,16	Term	User-defined signal 1
SIGUSR2	31,12,17	Term	User-defined signal 2
SIGCHLD	20,17,18	Ign	Child stopped or terminated
SIGCONT	19,18,25	Cont	Continue if stopped
SIGSTOP	17,19,23	Stop	Stop process
SIGTSTP	18,20,24	Stop	Stop typed at terminal
SIGTTIN	21,21,26	Stop	Terminal input for background process
SIGTTOU	22,22,27	Stop	Terminal output for background process

The signals SIGKILL and SIGSTOP cannot be caught, blocked, or ignored.

Next the signals not in the POSIX.1-1990 standard but described in SUSv2 and POSIX.1-2001.

Signal	Value	Action	Comment
SIGBUS	10,7,10	Core	Bus error (bad memory access)
SIGPOLL		Term	Pollable event (Sys V).

Manual page signal(7) line 122 (press h for help or q to quit)

With `man alarm`:

```

ALARM(2)                                Linux Programmer's Manual                                ALARM(2)

NAME
    alarm - set an alarm clock for delivery of a signal

SYNOPSIS
    #include <unistd.h>

    unsigned int alarm(unsigned int seconds);

DESCRIPTION
    alarm() arranges for a SIGALRM signal to be delivered to the calling process in seconds
    seconds.

    If seconds is zero, any pending alarm is canceled.

    In any event any previously set alarm() is canceled.

RETURN VALUE
    alarm() returns the number of seconds remaining until any previously scheduled alarm was
    due to be delivered, or zero if there was no previously scheduled alarm.

CONFORMING TO
    SVr4, POSIX.1-2001, 4.3BSD.

NOTES
    alarm() and setitimer(2) share the same timer; calls to one will interfere with use of the
    other.

    Alarms created by alarm() are preserved across execve(2) and are not inherited by children
    created via fork(2).

    sleep(3) may be implemented using SIGALRM; mixing calls to alarm() and sleep(3) is a bad
    idea.

    Scheduling delays can, as ever, cause the execution of the process to be delayed by an
    Manual page alarm(2) line 1 (press h for help or q to quit)

```

```

NOTES
    alarm() and setitimer(2) share the same timer; calls to one will interfere with use of the
    other.

    Alarms created by alarm() are preserved across execve(2) and are not inherited by children
    created via fork(2).

    sleep(3) may be implemented using SIGALRM; mixing calls to alarm() and sleep(3) is a bad
    idea.

    Scheduling delays can, as ever, cause the execution of the process to be delayed by an
    arbitrary amount of time.

SEE ALSO
    gettimeofday(2), pause(2), select(2), setitimer(2), sigaction(2), signal(2), sleep(3),
    time(7)

COLOPHON
    This page is part of release 3.74 of the Linux man-pages project. A description of the
    project, information about reporting bugs, and the latest version of this page, can be
    found at http://www.kernel.org/doc/man-pages/.

Linux                                2014-02-23                                ALARM(2)
Manual page alarm(2) line 15/52 (END) (press h for help or q to quit)

```

☒ Investigate the given code

☒ [FILL HERE with your description.]

This program will terminate itself in 5 seconds. At first, it will setup an alarm to be fired up in 5 seconds. Then it will print Looping forever ... and then it will enter an infinite loop. While the loop is executing, the SIGALRM will be delivered to program. This causes the application to be terminated because we don't have any signal handler for it.



[FILL HERE with screenshot of program execution]

```
alarm (5);
printf ("Looping forever . . . \n");
while (1);
printf("This line should never be executed\n");
return 0;
}
```

[ Wrote 9 lines ]

```
root@debian:~# make signalarm
cc      signalarm.c  -o signalarm
root@debian:~# ./signalarm
Looping forever . . .
Alarm clock
root@debian:~# _
```



Modify the given program by handling SIGALRM



[FILL HERE with your source code.]

The code is here:

```
GNU nano 2.2.6                               File: signalarm.c

#include <stdio.h>
#include <unistd.h>
#include <signal.h>

void signal_handler(int sig){
    printf("Signal number: %d", sig);
    printf("\n");
}

int main() {
    alarm (5);
    signal(SIGALRM, signal_handler);
    printf ("Looping forever . . . \n");
    pause();
    printf("This line should never be executed\n");
    return 0;
}
```

The result:

```
root@debian:~# cat signalarm.c
#include <stdio.h>
#include <unistd.h>
#include <signal.h>

void signal_handler(int sig){
    printf("Signal number: %d", sig);
    printf("\n");
}

int main() {
    alarm (5);
    signal(SIGALRM, signal_handler);
    printf ("Looping forever . . . \n");
    pause();
    printf("This line should never be executed\n");
    return 0;
}
root@debian:~# make signalarm
cc      signalarm.c  -o signalarm
root@debian:~# ./signalarm
Looping forever . . .
Signal number: 14
This line should never be executed
root@debian:~#
```

☒ Write a program that handles Ctrl + C

☒ [FILL HERE with your source code.]

The code is here:

GNU nano 2.2.6 File: double\_c.c

```
#include <stdio.h>
#include <signal.h>
#include <unistd.h>

void signal_handler(int signal_number) {
    puts("\nDo it again!");
}

int main() {
    puts("Press ctrl + c!");
    signal(SIGINT, signal_handler);
    pause();
    signal(SIGINT, NULL);
    pause();
    return 0;
}
```

The result:



```
root@debian:~# cat double_c.c
#include <stdio.h>
#include <signal.h>
#include <unistd.h>

void signal_handler(int signal_number) {
    puts("\nDo it again!");
}

int main() {
    puts("Press ctrl + c!");
    signal(SIGINT, signal_handler);
    pause();
    signal(SIGINT, NULL);
    pause();
    return 0;
}
root@debian:~# make double_c
cc double_c.c -o double_c
root@debian:~# ./double_c
Press ctrl + c!
^C
Do it again!
^C
root@debian:~#
```



Amirreza81 added the **documentation** label 5 days ago

  Amirreza81 assigned AMshoka 5 days ago

Assignees

 AMshoka

Labels

documentation

Projects

None yet

Milestone


No milestone

Development

Create a branch for this issue or link a pull request.

2 participants



 Pin issue 