



Sharif-OS-Lab /
summer1403-6-99101087_99100422 🗨

[Code](#)[Issues](#) **1**[Pull requests](#)[Actions](#)[Projects](#)[Security](#)[Insights](#)[Edit](#)[New issue](#)[Jump to bottom](#)

Session 6 Report #1

[Open](#)[36 tasks done](#)

BozorgmehrZia opened this issue on Aug 4 · 0 comments

Assignees



Labels

[documentation](#)

BozorgmehrZia commented on Aug 4 • edited ▾

Team Name: 99101087-99100422

Student Name of member 1: AmirReza Azari

Student No. of member 1: 99101087

Student Name of member 2: Bozorgmehr Zia

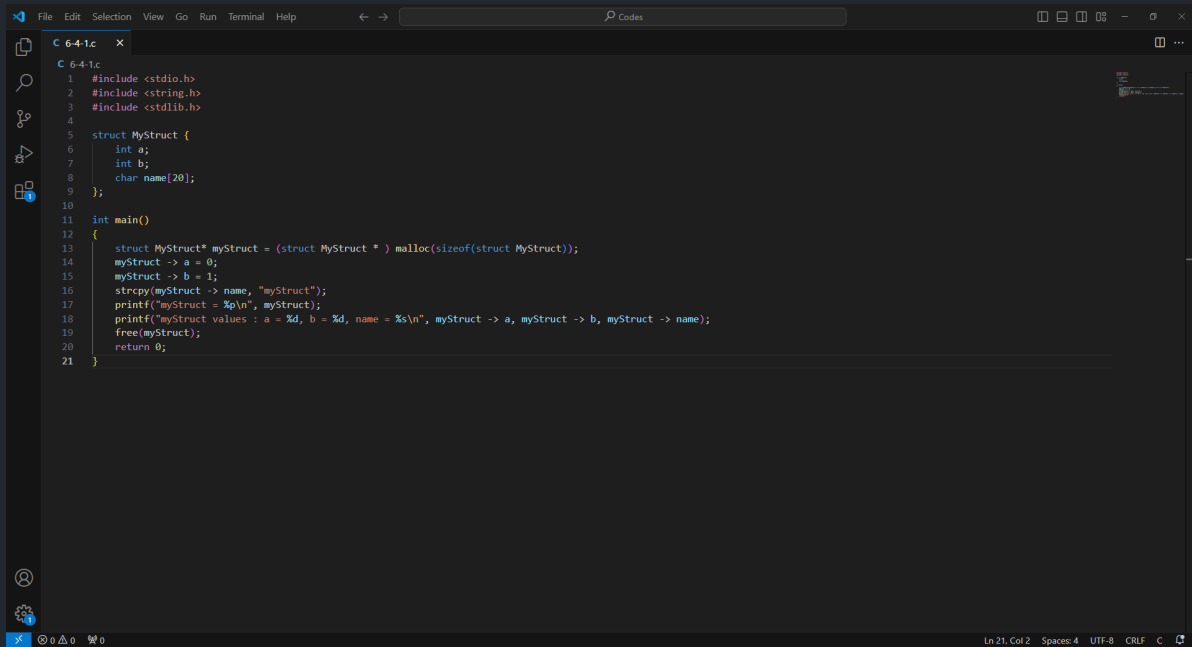
Student No. of member 2: 99100422

☒ Read Session Contents.

Section 6.4

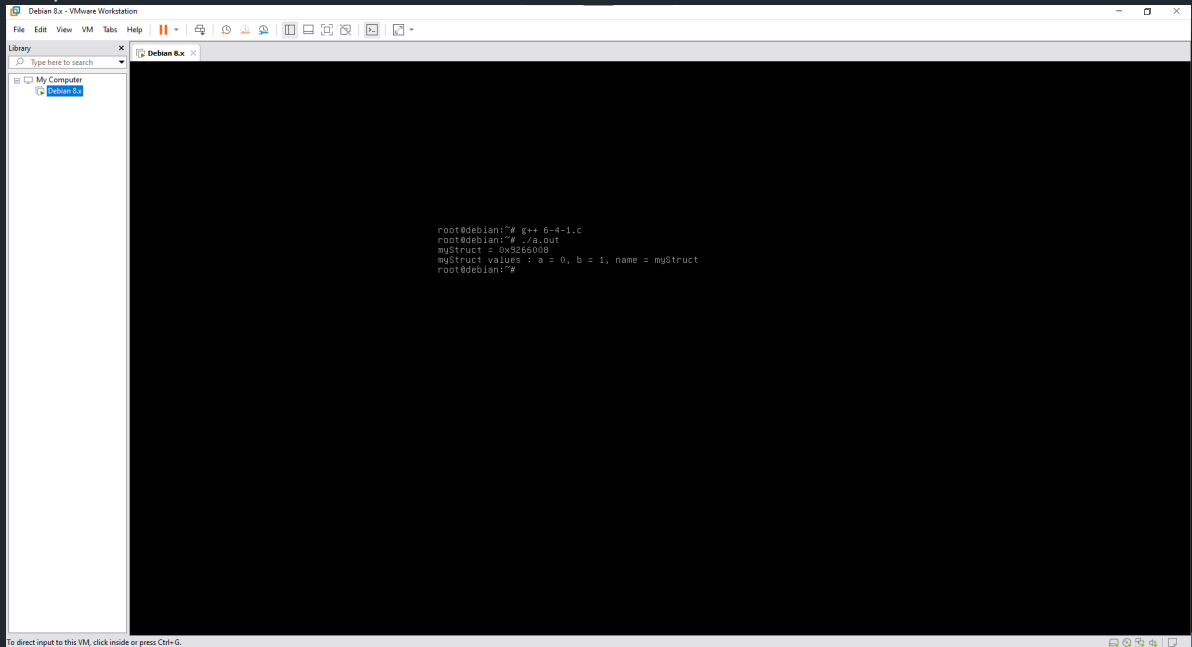
☒ Using `malloc` and `free` in program

✓ Code (in file 6-4-1.c):



```
1 #include <stdio.h>
2 #include <string.h>
3 #include <stdlib.h>
4
5 struct MyStruct {
6     int a;
7     int b;
8     char name[20];
9 };
10
11 int main()
12 {
13     struct MyStruct* myStruct = (struct MyStruct *) malloc(sizeof(struct MyStruct));
14     myStruct -> a = 0;
15     myStruct -> b = 1;
16     strcpy(myStruct -> name, "myStruct");
17     printf("myStruct = %p\n", myStruct);
18     printf("myStruct values : a = %d, b = %d, name = %s\n", myStruct -> a, myStruct -> b, myStruct -> name);
19     free(myStruct);
20     return 0;
21 }
```

✓ Output:



```
root@debian1:~# ./6-4-1.c
root@debian1:~# ./a.out
myStruct = 0x9266008
myStruct values : a = 0, b = 1, name = myStruct
root@debian1:~#
```

The `malloc` function returns a void pointer to the reserved space. The return value is NULL if not enough storage is available, or if size was specified as zero.

✓ Using `ps`



Debian 8.x - VMware Workstation

File Edit View VM Tabs Help

Library

Type here to search

My Computer

Debian 8.x

```
root@debian:~# ps -o user,vsz,rss,rss,rss,pmem,frame --e
USER      VSZ  RSS  PMEM  FRAME
root      5388 3864  0.1  system
root      0    0    0.0  kthreadd
root      0    0    0.0  ksoftirqd
root      0    0    0.0  kworker/
root      0    0    0.0  kworker/
root      0    0    0.0  kworker/
root      0    0    0.0  rcu_sche
root      0    0    0.0  rcu_bh
root      0    0    0.0  irqpoll
root      0    0    0.0  watchdog
root      0    0    0.0  khelper
root      0    0    0.0  kdsystm
root      0    0    0.0  netns
root      0    0    0.0  khungtad
root      0    0    0.0  amibac
root      0    0    0.0  ksm
root      0    0    0.0  khugepag
root      0    0    0.0  crpd10
root      0    0    0.0  krtstr
root      0    0    0.0  bioset
root      0    0    0.0  kblockd
root      0    0    0.0  kworker/
```

To direct input to this VM, click inside or press Ctrl+G.



Debian 8.x - VMware Workstation

File Edit View VM Tabs Help

Library


Type here to search

My Computer

Debian 8.x

```
root      0    0    0.0  kworker/
root      0    0    0.0  ksm
root      0    0    0.0  vmstat
root      0    0    0.0  fsnotif
root      0    0    0.0  kthreadd
root      0    0    0.0  inv6.add
root      0    0    0.0  deferreq
root      0    0    0.0  kworker/
root      0    0    0.0  kworker/
root      0    0    0.0  khudd
root      0    0    0.0  ata_sff
root      0    0    0.0  mpt_poll
root      0    0    0.0  mpt0
root      0    0    0.0  kpsmouse
root      0    0    0.0  scsi-sh
root      0    0    0.0  scsi-lm
root      0    0    0.0  kworker/
root      0    0    0.0  scsi-sh
root      0    0    0.0  scsi-lm
root      0    0    0.0  scsi-sh
root      0    0    0.0  scsi-lm
root      0    0    0.0  kworker/
root      0    0    0.0  kworker/
root      0    0    0.0  jbd2/sda
root      0    0    0.0  ext4-rsv
```



To direct input to this VM, click inside or press Ctrl+G.



Columns are:

- ✓ **user** : This column displays the user ID or username of the person who owns and runs the process. Example Output: root, user1. Alias: uname.
- ✓ **vsz** : This column shows the virtual memory size of the process in kilobytes. It's the total memory size that a process can access, such as the memory of the shared libraries, the allocated memory such as stack and heap, and swapped out memory. Example Output: 1048576 (for 1 GB). Alias: vsize.
- ✓ **rss** : This column shows the resident set size, which is the non-swapped physical memory a task has used (in kilobytes). It doesn't show all of the information about the memory, like loaded libraries, heap, or stack. It shows the shared libraries if their pages are in memory. Example Output: 2048 (for 2 MB).
- ✓ **pmem** : This column shows the percentage of resident set size (RSS) to the physical memory used by the process. Example Output: 0.1
- ✓ **fname** : This column displays the first 8 bytes of the process's executable file name. The output in this column may contain spaces. Example Output: bash, python

Getting started with memory segments

**which ls****size /bin/ls**

As you see, 116455 bytes from memory are allocated to text part, 1140 bytes for initialized data, 3252 bytes for bss, and 120847 bytes for dec.

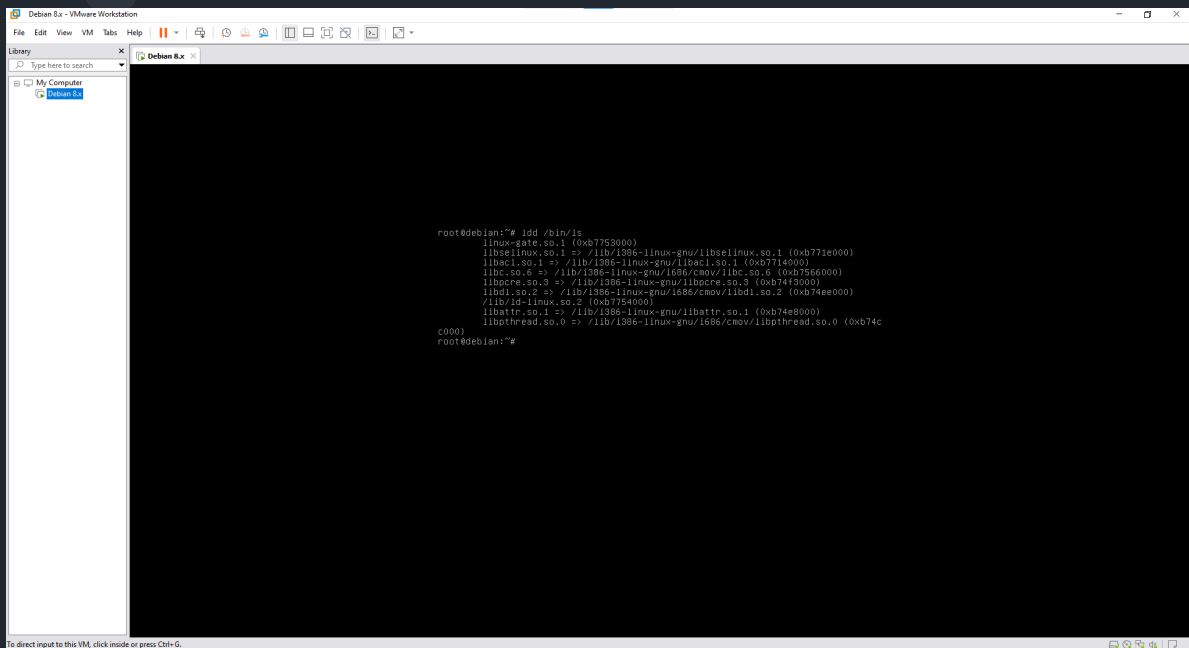


This command doesn't show stack and heap size.



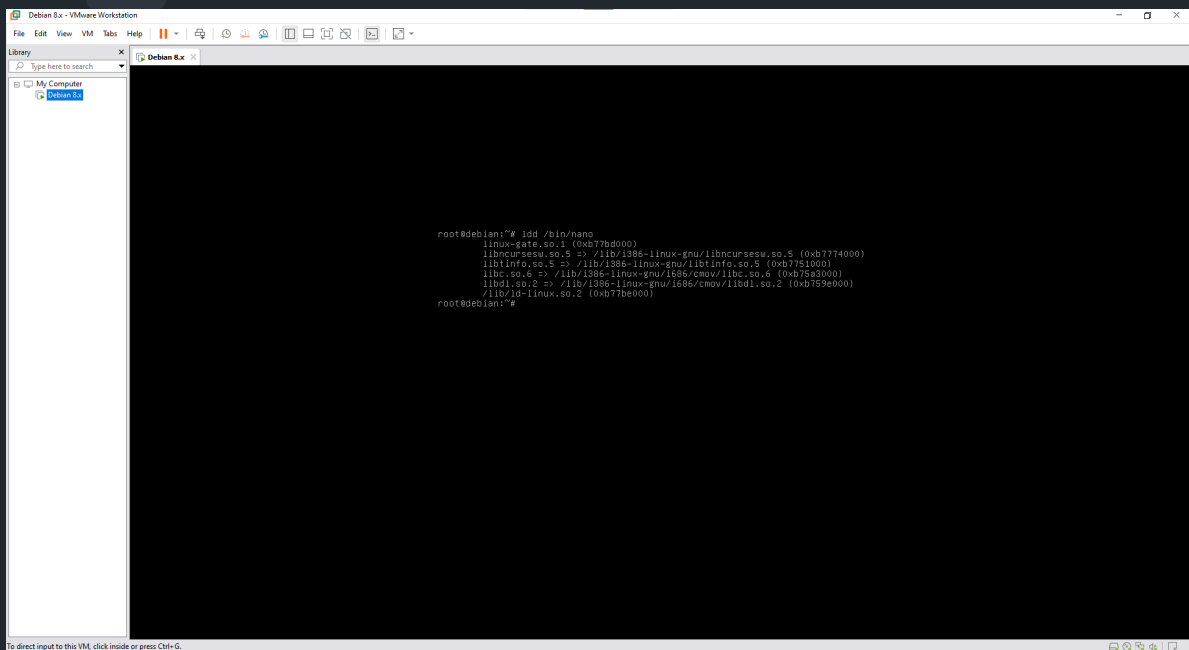
Getting started with memory sharing

✓ lld for ls



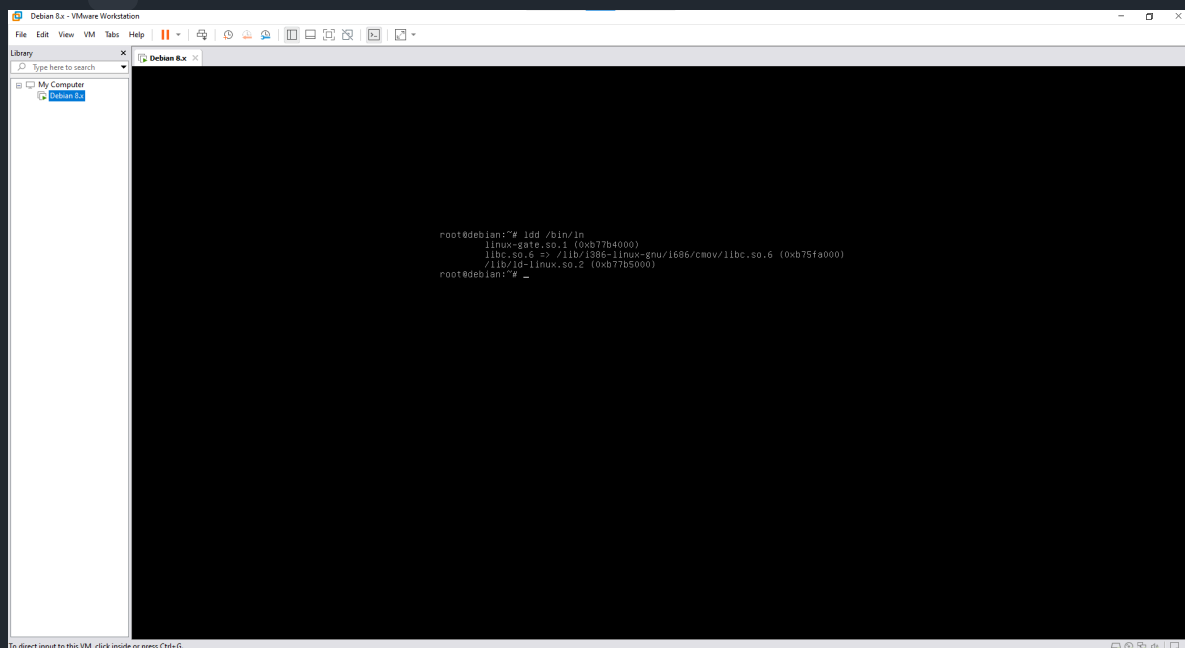
```
root@debian:~# ldd /bin/ls
linux-gate.so.1 (0xb7753000)
libselinux.so.1 => /lib/1986-linux-gnu/libselinux.so.1 (0xb771e000)
libsc1.so.1 => /lib/1986-linux-gnu/libsc1.so.1 (0xb7714000)
libc.so.6 => /lib/1986-linux-gnu/libc.so.6 (0xb7560000)
libpcre.so.3 => /lib/1986-linux-gnu/libpcre.so.3 (0xb74f3000)
libdl.so.2 => /lib/1986-linux-gnu/libdl.so.2 (0xb74ee000)
/lib/ld-linux.so.2 (0xb7754000)
libattr.so.1 => /lib/1986-linux-gnu/libattr.so.1 (0xb74e0000)
libpthread.so.0 => /lib/1986-linux-gnu/libpthread.so.0 (0xb74c0000)
root@debian:~#
```

✓ lld for nano



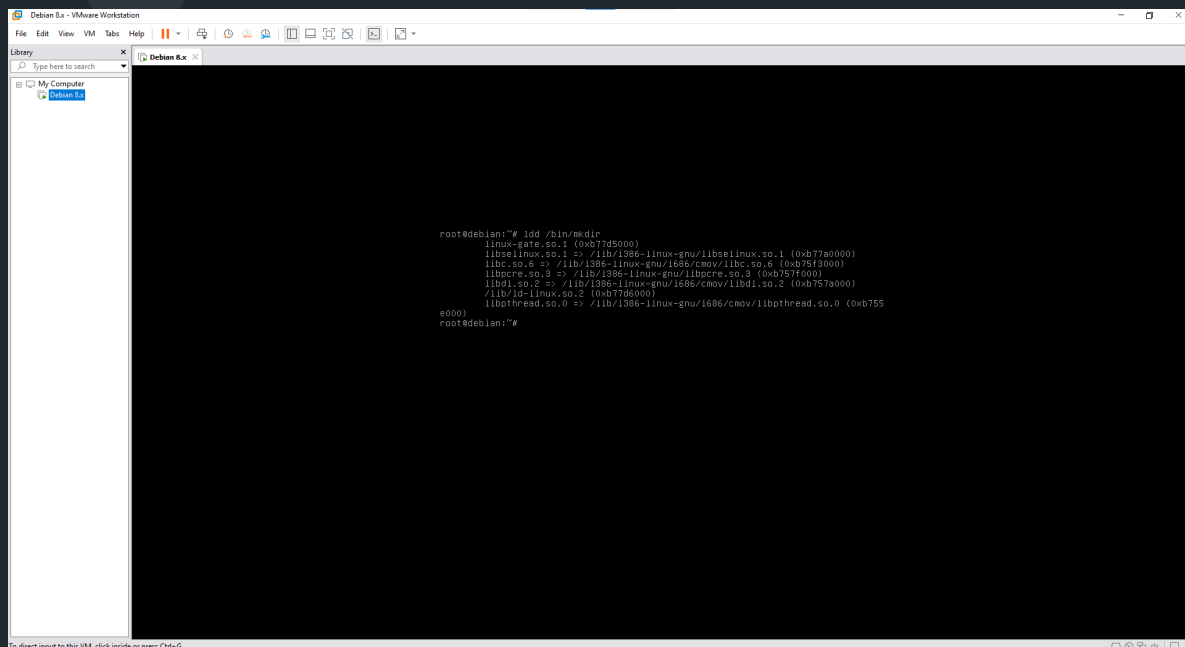
```
root@debian:~# ldd /bin/nano
linux-gate.so.1 (0xb77b0000)
libncursesw.so.5 => /lib/1986-linux-gnu/libncursesw.so.5 (0xb7740000)
libtinfo.so.5 => /lib/1986-linux-gnu/libtinfo.so.5 (0xb7751000)
libc.so.6 => /lib/1986-linux-gnu/libc.so.6 (0xb75a3000)
libdl.so.2 => /lib/1986-linux-gnu/libdl.so.2 (0xb759e000)
/lib/ld-linux.so.2 (0xb77be000)
root@debian:~#
```

✓ lld for ln



```
root@debian:~# lld /bin/ln
linux-gate.so.1 (0xb77b4000)
libc.so.6 => /lib/i386-linux-gnu/i686/cmov/libc.so.6 (0xb75fa000)
/lib/ld-linux.so.2 (0xb7705000)
root@debian:~#
```

✓ lld for mkdir



```
root@debian:~# lld /bin/mkdir
linux-gate.so.1 (0xb77d5000)
libselinux.so.1 => /lib/i386-linux-gnu/libselinux.so.1 (0xb77a0000)
libc.so.6 => /lib/i386-linux-gnu/i686/cmov/libc.so.6 (0xb75fa000)
libpcre.so.3 => /lib/i386-linux-gnu/libpcre.so.3 (0xb757f000)
libdl.so.2 => /lib/i386-linux-gnu/i686/cmov/libdl.so.2 (0xb757a000)
/lib/ld-linux.so.2 (0xb77d5000)
libpthread.so.0 => /lib/i386-linux-gnu/i686/cmov/libpthread.so.0 (0xb755
8000)
root@debian:~#
```

✓ Getting started with addresses

man etext

```
Debian 8.x - VMware Workstation
File Edit View VM Tabs Help
Library
Type here to search
My Computer
Debian 8.x

END(3) Linux Programmer's Manual END(3)

NAME
etext, edata, end - end of program segments

SYNOPSIS
extern etext;
extern edata;
extern end;

DESCRIPTION
The addresses of these symbols indicate the end of various program segments:

etext This is the first address past the end of the text segment (the program code).

edata This is the first address past the end of the initialized data segment.

end This is the first address past the end of the uninitialized data segment (also known as the BSS segment).

CONFORMING TO
Manus: pose etext(3), line 1 (press h for help or q to quit)
```

```
Debian 8.x - VMware Workstation
File Edit View VM Tabs Help
Library
Type here to search
My Computer
Debian 8.x

CONFORMING TO
although these symbols have long been provided on most UNIX systems, they are not standardized; use with caution.

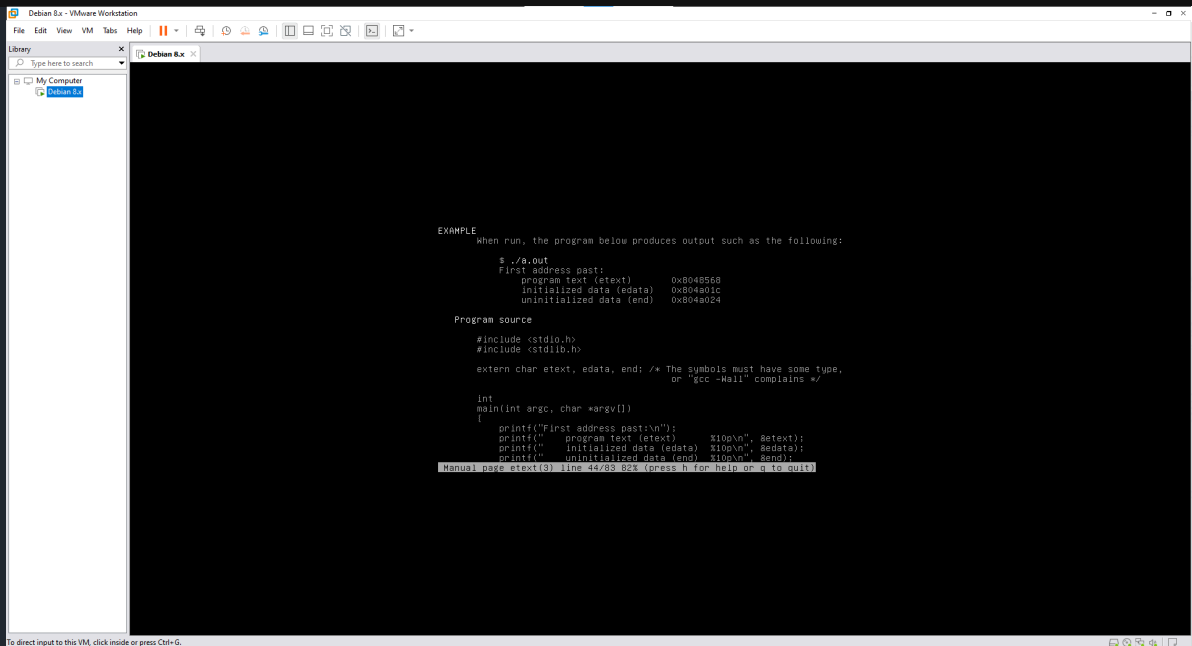
NOTES
The program must explicitly declare these symbols; they are not defined in any header file.

On some systems the names of these symbols are preceded by underscores, thus: _etext, _edata, and _end. These symbols are also defined for programs compiled on Linux.

At the start of program execution, the program break will be somewhere near &end (perhaps at the start of the following page). However, the break will change as memory is allocated via brk(2) or malloc(3). Use sbrk(2) with an argument of zero to find the current value of the program break.

EXAMPLE
When run, the program below produces output such as the following:

$ ./a.out
First address past:
program text (etext) 0x8048569
Manus: pose etext(3), line 26/88 61s (press h for help or q to quit)
```

Debian 8.x - VMware Workstation

Library

Type here to search

My Computer

Debian 8.x

```
EXAMPLE
When run, the program below produces output such as the following:

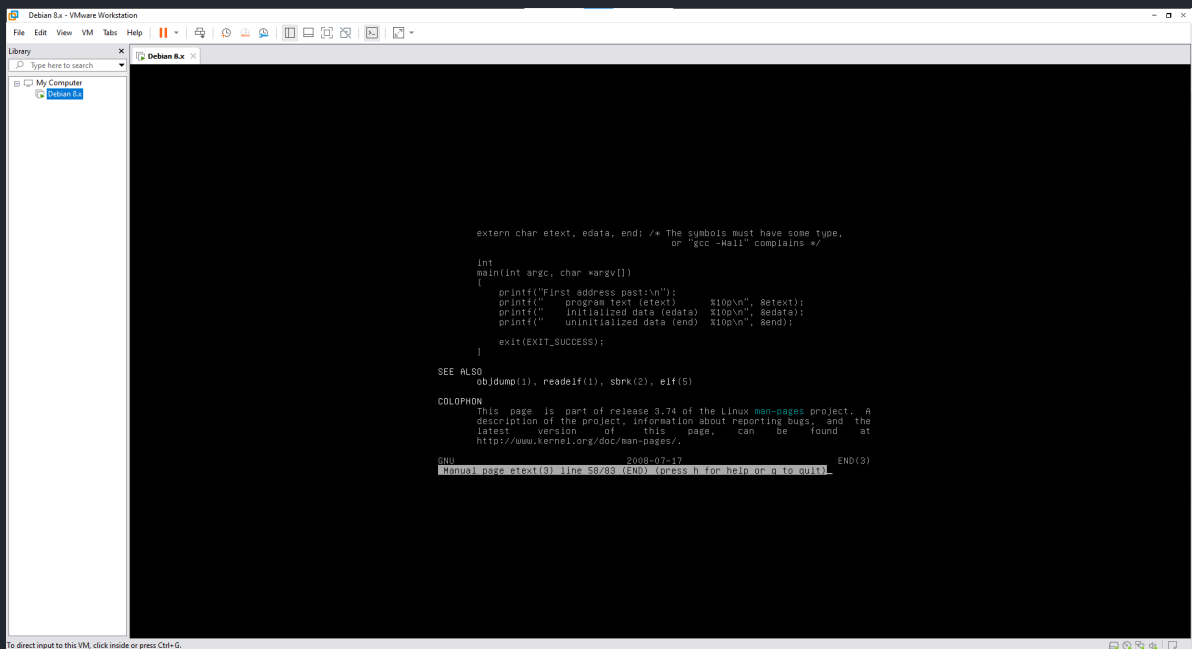
$ ./a.out
First address past:
  program text (etext)      0x8048568
  initialized data (edata)  0x804a01c
  uninitialized data (end)  0x804a024

Program source
#include <stdio.h>
#include <stdlib.h>

extern char etext, edata, end; /* The symbols must have some type,
                               or 'gcc -Wall' complains */

int
main(int argc, char *argv[])
{
    printf("First address past:\n");
    printf("  program text (etext)      %10p\n", &etext);
    printf("  initialized data (edata)    %10p\n", &edata);
    printf("  uninitialized data (end)    %10p\n", &end);
Manual page etext(3) line 44/83 8c (press h for help or q to quit)
```

To direct input to this VM, click inside or press Ctrl+G.



Debian 8.x - VMware Workstation

Library

Type here to search

My Computer

Debian 8.x

```
extern char etext, edata, end; /* The symbols must have some type,
                               or 'gcc -Wall' complains */

int
main(int argc, char *argv[])
{
    printf("First address past:\n");
    printf("  program text (etext)      %10p\n", &etext);
    printf("  initialized data (edata)    %10p\n", &edata);
    printf("  uninitialized data (end)    %10p\n", &end);
    exit(EXIT_SUCCESS);
}

SEE ALSO
objdump(1), readelf(1), sbrk(2), elf(5)

COLOPHON
This page is part of release 3.74 of the Linux man-pages project. A
description of the project, information about reporting bugs, and the
latest version of this page, can be found at
http://www.kernel.org/doc/man-pages/.

GNU          2008-07-17          END(3)
Manual page etext(3) line 58/83 (END) (press h for help or q to quit)
```

To direct input to this VM, click inside or press Ctrl+G.

✓ Code (in file 6-4-2.c):

```

1 #include <stdio.h>
2 #include <stdlib.h>
3
4 extern char etext, edata, end; /* The symbols must have some type,
5                               or "gcc -Wall" complains */
6
7 int main(int argc, char *argv[])
8 {
9     printf("First address past:\n");
10    printf("    program text (etext)    %10p\n", &etext);
11    printf("    initialized data (edata) %10p\n", &edata);
12    printf("    uninitialized data (end) %10p\n", &end);
13
14    exit(EXIT_SUCCESS);
15 }
16

```

Output:

```

root@debian:~# g++ 6-4-2.c
root@debian:~# ./a.out
First address past:
    program text (etext)    0x80405f8
    initialized data (edata) 0x8040898
    uninitialized data (end) 0x804089c
root@debian:~# _

```

As you see, the addresses are exactly the same as what we expected.

✓ Description about `etext` , the meaning of the comment, why using the term symbol, why using `extern`:

✓ The symbols are:

- ✓ `etext` : This is the first address past the end of the text segment (the program code). The text segment contains the compiled machine code of the program.
- ✓ `edata` : This is the first address past the end of the initialized data segment
- ✓ `end` : This is the first address past the end of the uninitialized data segment (also known as the BSS segment).

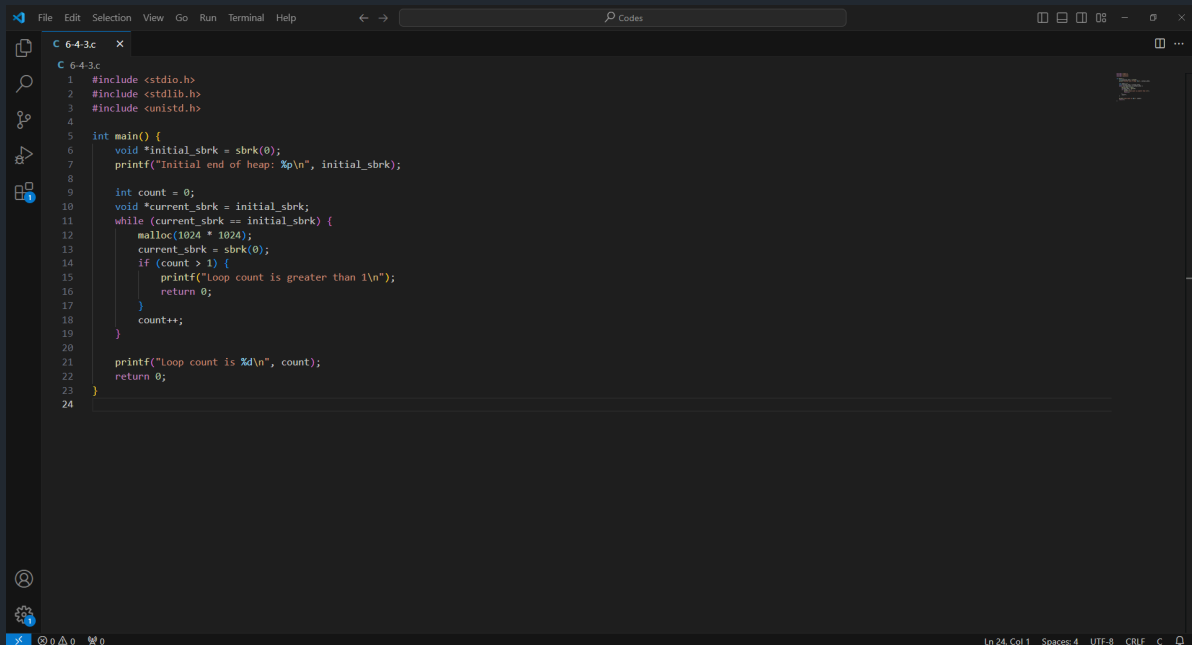
- ✓ Meaning of the comment: The '-Wall' flag enables all compiler's warning messages, which helps in writing cleaner and more error-free code. In C, declaring symbols with `extern`

requires specifying a type. Without a type, the compiler would issue warnings or errors. Here, char is chosen as the type because it is a basic type that simply gives these symbols a valid type without any specific size or alignment constraints.

- ✓ Why use the term symbol: These symbols () are used to refer to specific locations in memory and do not hold values themselves. They are part of the program's memory map, which is why they are called symbols rather than variables.
- ✓ Why use extern: The extern keyword is used to declare these symbols without defining them. This informs the compiler that these symbols are defined elsewhere, typically by the linker. These symbols are defined by the linker, not by the C code, so using extern tells the compiler that their actual addresses will be resolved during the linking stage. Also, by declaring them as extern char, the code can take the address of these symbols (using &etext, &edata, &end) to print or otherwise.


✓ sbrk analysis:

Code (in file 6-4-3.c):



```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 int main() {
6     void *initial_sbrk = sbrk(0);
7     printf("Initial end of heap: %p\n", initial_sbrk);
8
9     int count = 0;
10    void *current_sbrk = initial_sbrk;
11    while (current_sbrk == initial_sbrk) {
12        malloc(1024 * 1024);
13        current_sbrk = sbrk(0);
14        if (count > 1) {
15            printf("Loop count is greater than 1\n");
16            return 0;
17        }
18        count++;
19    }
20
21    printf("Loop count is %d\n", count);
22    return 0;
23 }
24
```

Output:



As you see, the loop count is not equal to 1. That's because memory allocators often request larger chunks of memory from the operating system and then manage these chunks internally. When malloc is called, it might be allocating memory from a pre-existing pool rather than extending the heap. This means the end of the heap doesn't change with every call to malloc. In the above code, the size of malloc was 1 MB. If we change it to 1 KB, the address of end of heap will change and the loop count will be 1. Its output will be:

We could write code like this to see if malloc changes the address of end of heap:

```


#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>


int main() {
    printf("End of heap before malloc: %p\n", sbrk(0));
    malloc(1024 * 1024);
    printf("End of heap after malloc: %p\n", sbrk(0));
    return 0;
}

```

- ☒ stack growth analysis:
- Code (in file 6-4-4.c):

Output:





As you see, addresses decrease because the stack grows backward (from higher address to lower address).

The codes are:

[Codes.zip](#)



Bozorgmehrzia added the **documentation** label on Aug 4



Bozorgmehrzia assigned AMshoka on Aug 4

Assignees



AMshoka

Labels



documentation

Projects



None yet

Milestone



No milestone


Development



Create a [branch](#) for this issue or link a pull request.

2 participants



 Pin issue ⓘ